

CENG 790: Big Data Analytics, Fall 2020

Report of Assignment 3: Random Forest Classifier

Extract Features

To build a classifier model, you first extract the features that most contribute to the classification. In the credit data set the data is labeled with two classes – 1 (**creditable**) and 0 (**not creditable**).

The features for each item consist of the fields shown below:

- Label → **creditable**: 0 or 1
- Features → {"balance", "duration", "history", "purpose", "amount", "savings", "employment", "instPercent", "sexMarried", "guarantors", "residenceDuration", "assets", "age", "concCredit", "apartment", "credits", "occupation", "dependents", "hasPhone", "foreign"}

In order for the features to be used by a machine learning algorithm, the features are transformed and put into Feature Vectors, which are vectors of numbers representing the value for each feature.

1. Use a **VectorAssembler** to transform and return a new dataframe with all of the feature columns in a vector column.

```
56 // EXTRACT FEATURES
57 // 1. Use a VectorAssembler to transform and return a new dataframe with all of the feature columns in a vector
58 // column.
59 val featureColumnNames = creditDF.columns.filter(colName => !colName.equals("creditability"))
60 val featureColumnName = "features"
61
62 val featuresAssembler = new VectorAssembler()
63   .setInputCols(featureColumnNames)
64   .setOutputCol(featureColumnName)
```

2. Use a **StringIndexer** to return a Dataframe with the creditability column added as a label.

```
66 // 2. Use a StringIndexer to return a Dataframe with the creditability column added as a label
67 val labelColumnName = "label"
68
69 val labelIndexer = new StringIndexer()
70   .setInputCol("creditability")
71   .setOutputCol(labelColumnName)
72   .fit(creditDF)
```

3. Use **randomSplit** function to split the data into two sets: 75% of the data is used to train (and tune) the model, 25% will be used for testing.

```
75 // 3. Use randomSplit function to split the data into two sets: 75% of the data is used to train (and tune) the
76 // model, 25% will be used for testing.
77 val Array(trainCreditDF, testCreditDF) = creditDF
78   .randomSplit(Array(0.75, 0.25), seed = 4321)
```

Train the model and optimize hyperparameters

The model is trained by making associations between the input features and the labeled output associated with those features. In order to find the best model, we search for the optimal combinations of the classifier parameters.

You will optimize the model using a pipeline. A pipeline provides a simple way to try out different combinations of parameters, using a process called grid search, where you set up the parameters to test, and MLLib will test all the combinations. **Pipelines** make it easy to tune an entire model building workflow at once, rather than tuning each element in the Pipeline separately.

4. Next, we train a RandomForestClassifier with the parameters:

- a. **maxDepth**: Maximum depth of a tree. Increasing the depth makes the model more powerful, but deep trees take longer to train.
- b. **maxBins**: Maximum number of bins used for discretizing continuous features and for choosing how to split on features at each node.
- c. **impurity**: Criterion used for information gain calculation
- d. **auto**: Automatically select the number of features to consider for splits at each tree node
- e. **seed**: Use a random seed number, allowing to repeat the results. Use the random seed 4321 in this assignment.

```
80 // TRAIN THE MODEL AND OPTIMIZE HYPERPARAMETERS
81 // 4. Next, we train a RandomForest Classifier with the parameters:
82 //   a. maxDepth: Maximum depth of a tree. Increasing the depth makes the model more powerful, but deep trees take
83 //      longer to train.
84 //   b. maxBins: Maximum number of bins used for discretizing continuous features and for choosing how to split on
85 //      features at each node.
86 //   c. impurity: Criterion used for information gain calculation
87 //   d. auto: Automatically select the number of features to consider for splits at each tree node
88 //   e. seed: Use a random seed number, allowing to repeat the results. Use the random seed 4321 in this
89 //      assignment.
90 // The model is trained by making associations between the input features and the labeled output associated with
91 // those features. In order to find the best model, we search for the optimal combinations of the classifier
92 // parameters.
93 // You will optimize the model using a pipeline. A pipeline provides a simple way to try out different combinations
94 // of parameters, using a process called grid search, where you set up the parameters to test, and MLLib will test
95 // all the combinations. Pipelines make it easy to tune an entire model building workflow at once, rather than
96 // tuning each element in the Pipeline separately.
97 val randomForestClassifier = new RandomForestClassifier()
98   .setFeaturesCol(featureColumnName)
99   .setLabelCol(labelColumnName)
100   .setSeed(4321)
```

Use the **ParamGridBuilder** utility to construct the parameter grid with the following values: maxBins [24, 28, 32], maxDepth, [3, 5, 7], impurity ["entropy", "gini"]

```
102 // Use the ParamGridBuilder utility to construct the parameter grid with the following values: maxBins [24, 28, 32],
103 // maxDepth, [3, 5, 7], impurity ["entropy", "gini"]
104 val paramGridBuilder = new ParamGridBuilder()
105   .addGrid(randomForestClassifier.featureSubsetStrategy, Array("auto"))
106   .addGrid(randomForestClassifier.impurity, Array("entropy", "gini"))
107   .addGrid(randomForestClassifier.maxBins, Array(24, 28, 32))
108   .addGrid(randomForestClassifier.maxDepth, Array(3, 5, 7))
109   .build()
```

Also, I tried other parameters such as **numTrees**, **subsamplingRate** and different values for **featureSubsetStrategy**, **maxBins**, **maxDepth** in the tuning phase. I will mention about it later.

5. Next, you will create and set up a pipeline to make things easier. A Pipeline consists of a sequence of stages, each of which is either an Estimator or a Transformer.

Use **TrainValidationSplit** that creates a (training, test) dataset pair. It splits the dataset into these two

parts using the `trainRatio` parameter. For example, with `trainRatio=0.75`, `TrainValidationSplit` will generate a training and test dataset pair where 75% of the data is used for training and 25% for validation. Use these values in your code. **Note that, this is different from the original random data split.** Here, we further divide the training dataset into training and validation set, for tuning purposes. The final model that has been tuned will be used to evaluate the result on the test set.

The `TrainValidationSplit` uses an Estimator, a set of `ParamMaps`, and an Evaluator. Estimator should be your random forest model, the `ParamMaps` is the parameter grid that you built in the previous step. The Evaluator should be **`new BinaryClassificationEvaluator()`**.

```

122 // 5. Next, you will create and set up a pipeline to make things easier. A Pipeline consists of a sequence of
123 // stages, each of which is either an Estimator or a Transformer.
124 // Use TrainValidationSplit that creates a (training, test) dataset pair. It splits the dataset into these two parts
125 // using the trainRatio parameter. For example, with trainRatio=0.75, TrainValidationSplit will generate a training
126 // and test dataset pair where 75% of the data is used for training and 25% for validation. Use these values in your
127 // code. Note that, this is different from the original random data split. Here, we further divide the training
128 // dataset into training and validation set, for tuning purposes. The final model that has been tuned will be used
129 // to evaluate the result on the test set.
130 val binaryClassificationEvaluator = new BinaryClassificationEvaluator()
131   .setLabelCol(labelColumnName)

133 // The TrainValidationSplit uses an Estimator, a set of ParamMaps, and an Evaluator. Estimator should be your random
134 // forest model, the ParamMaps is the parameter grid that you built in the previous step. The Evaluator should be
135 // new BinaryClassificationEvaluator().
136 val trainValidationSplit = new TrainValidationSplit()
137   .setEstimator(randomForestClassifier)
138   .setEstimatorParamMaps(paramGridBuilder)
139   .setEvaluator(binaryClassificationEvaluator)
140   .setTrainRatio(0.75)
141   .setSeed(4321)

143 val pipeline = new Pipeline()
144   .setStages(Array(featuresAssembler, labelIndexer, trainValidationSplit))

146 val model = pipeline.fit(trainCreditDF)
147 model.write.overwrite().save(MODEL_PATH)

```

What are the best combinations for the hyper parameters you optimized?

```

149 val bestModel = model.stages(2).asInstanceOf[TrainValidationSplitModel]
150   .bestModel.asInstanceOf[RandomForestClassificationModel]
151 val impurity = bestModel.getImpurity
152 val maxBins = bestModel.getMaxBins
153 val maxDepth = bestModel.getMaxDepth
154 println(s"""Model's Parameters => Impurity:\$impurity\", MaxBins:\$maxBins, MaxDepth:\$maxDepth""")

Model's Parameters => Impurity:"entropy", MaxBins:32, MaxDepth:7

```

The combination for the best model:

- Impurity is **entropy**,
- MaxBins is **32**,
- MaxDepth is **7**.

6. Finally, evaluate the pipeline best-fitted model by comparing test predictions with test labels. You can use **`transform`** function to get the predictions for test dataset. You can use evaluator's **`evaluate`** function to get the metrics.

What are your accuracy values on train and test sets? Feel free to provide more stats and comment on your performance.

```

155 // 6. Finally, evaluate the pipeline best-fitted model by comparing test predictions with test labels. You can use
156 // transform function to get the predictions for test dataset. You can use evaluator's evaluate function to get the
157 // metrics.
158 val trainPredictions = model.transform(trainCreditDF)
159 val testPredictions = model.transform(testCreditDF)
160
161 val trainResult = binaryClassificationEvaluator.evaluate(trainPredictions)
162 val testResult = binaryClassificationEvaluator.evaluate(testPredictions)
163 println(s"Model's Accuracies on => Train Data:$trainResult, Test Data:$testResult")

```

Model's Accuracies on => Train Data:0.9767933538951346, Test Data:0.7447205977907733

The model's accuracy on train data is **0.9767933538951346**, and the accuracy on test data is **0.7447205977907733**. However, the accuracy on the test data didn't satisfy me as I had expected higher accuracy value. The reason for that might be the train and validation data size, which is 75% of all data (750 sample), was not enough or the tuning parameters were not good enough.

Therefore, I tried extending the parameter tuning phase in order to achieve better result:

```

110 // Additional ParamGridBuilders
111 val extraParamGridBuilder = new ParamGridBuilder()
112   .addGrid(randomForestClassifier.featureSubsetStrategy, Array("auto") ++ Array("all", "onethird", "sqrt", "log2"))
113   .addGrid(randomForestClassifier.impurity, Array("entropy", "gini"))
114   .addGrid(randomForestClassifier.maxBins, Array(24, 28, 32) ++ Array(12, 16, 20, 36, 40))
115   .addGrid(randomForestClassifier.maxDepth, Array(3, 5, 7) ++ Array(2, 4, 6, 8, 9, 10))
116   .addGrid(randomForestClassifier.numTrees, Array(20, 24, 28, 32))
117   .addGrid(randomForestClassifier.subsamplingRate, Array(0.1, 0.25, 0.5, 0.75, 1.0))
118   .build()

```

```

132 // The TrainValidationSplit uses an Estimator, a set of ParamMaps, and an Evaluator. Estimator should be your random
133 // forest model, the ParamMaps is the parameter grid that you built in the previous step. The Evaluator should be
134 // new BinaryClassificationEvaluator().
135 val trainValidationSplit = new TrainValidationSplit()
136   .setEstimator(randomForestClassifier)
137   .setEstimatorParamMaps(extraParamGridBuilder)
138   .setEvaluator(binaryClassificationEvaluator)
139   .setTrainRatio(0.75)
140   .setSeed(4321)

```

```

150 val bestModel = model.stages(2).asInstanceOf[TrainValidationSplitModel]
151   .bestModel.asInstanceOf[RandomForestClassificationModel]
152 //val bestModel = model.stages(2).asInstanceOf[CrossValidatorModel]
153 // .bestModel.asInstanceOf[RandomForestClassificationModel]
154 val impurity = bestModel.getImpurity
155 val maxBins = bestModel.getMaxBins
156 val maxDepth = bestModel.getMaxDepth
157 // println(s"Model's Parameters => Impurity:$impurity, MaxBins:$maxBins, MaxDepth:$maxDepth")
158 val featureSubsetStrategy = bestModel.getFeatureSubsetStrategy
159 val numTrees = bestModel.getNumTrees
160 val subsamplingRate = bestModel.getSubsamplingRate
161 println(s"Model's Parameters => FeatureSubsetStrategy:$featureSubsetStrategy, Impurity:$impurity, MaxBins:$maxBins, MaxDepth:$maxDepth, NumTrees:$numTrees, SubsamplingRate:$subsamplingRate")

```

Model's Parameters => FeatureSubsetStrategy:"onethird", Impurity:"entropy", MaxBins:36, MaxDepth:9, NumTrees:24, SubsamplingRate:0.75

Model's Accuracies on => Train Data:0.9909876297549167, Test Data:0.7828541260558806

- I extended the ParamGridBuilder with the following parameters:
 - FeatureSubsetStrategy -> **auto**, all, onethird, sqrt, log2
 - Impurity -> **entropy**, gini
 - MaxBins -> 12, 16, 20, **24, 28, 32**, 34, 36, 38, 40
 - MaxDepth -> 2, **3**, 4, **5**, 6, **7**, 8, 9, 10
 - NumTrees -> 4, 8, 12, 16, **20**, 24, 28, 32
 - SubsamplingRate -> 0.1, 0.25, 0.5, 0.75, **1.0**
- Normally, there was 18 different combinations. With the extended parameters, there was 36000 different combinations and I was sure that the model would be better after the learning process. It took around 2 hours to complete the model fitting phase. The best combination was:
 - FeatureSubsetStrategy is **onethird**,
 - Impurity is **entropy**,
 - MaxBins is **36**,

- MaxDepth is **9**,
- NumTrees is **24**,
- SubsamplingRate is **0.75**.
- The new best model achieved **0.9909876297549167** accuracy on training data and **0.7828541260558806** accuracy, which is better than the previous best model's accuracy, on test data.

In addition, I also tried changing randomSplit and trainValidationSplit ratios, but it didn't improve the model's accuracy neither. I even tried removing the seed with the thought of **seed 4321** was the unlucky one and running the training process multiple times in order to achieve a better result but that was not the case.

I also tried using **CrossValidator** with different number of folds such as 2, 4, 6, 8, 10 instead of TrainValidationSplit in order to achieve better score on test data. CrossValidator achieved **0.9748947036779096** on train data and **0.7854938271604937** on test data with the following setting: numFold **2**, impurity **entropy**, maxBins **24** and maxDepth **7**. CrossValidator's accuracy on test data was the highest so far.

```
141 // Try CrossValidator
142 val crossValidator = new CrossValidator()
143   .setEstimator(randomForestClassifier)
144   .setEstimatorParamMaps(paramGridBuilder)
145   .setEvaluator(binaryClassificationEvaluator)
146   .setNumFolds(10)
147   .setSeed(4321)

151 val pipeline = new Pipeline()
152   .setStages(Array(featuresAssembler, labelIndexer, crossValidator))

159 val bestModel = model.stages(2).asInstanceOf[CrossValidatorModel]
160   .bestModel.asInstanceOf[RandomForestClassificationModel]

Model's Parameters => Impurity:"entropy", MaxBins:24, MaxDepth:7
Model's Accuracies on => Train Data:0.9748947036779096, Test Data:0.7854938271604937
```

Lastly, I tried using CrossValidator with the following parameters (which are the combination of two best models): numFold **2**, featureSubsetStrategy **onethird**, impurity **entropy**, maxBins **36**, maxDepth **9**, numTrees **24**, subsamplingRate **0.75**.

It achieved accuracy **0.9909876297549167** on training data, **0.7828541260558807** on test data.

```
Model's Parameters => Impurity:"entropy", MaxBins:36, MaxDepth:9
Model's Accuracies on => Train Data:0.9909876297549167, Test Data:0.7828541260558807
```

To conclude, using TrainValidationSplit with featureSubsetStrategy **onethird**, impurity **entropy**, maxBins **36**, maxDepth **9**, numTrees **24**, and subsamplingRate **0.75** or using CrossValidator with numFold **2**, featureSubsetStrategy **onethird**, impurity **entropy**, maxBins **36**, maxDepth **9**, numTrees **24**, and subsamplingRate **0.75** achieved highest accuracy on training data (**0.991**) whereas using TrainValidationSplit with featureSubsetStrategy **onethird**, impurity **entropy**, maxBins **36**, maxDepth **9**, numTrees **24**, subsamplingRate **0.75** achieved highest accuracy on test data (**0.786**).

Maybe, extended parameter grid build can be run again with CrossValidator with different numFold parameters in order to find better model.